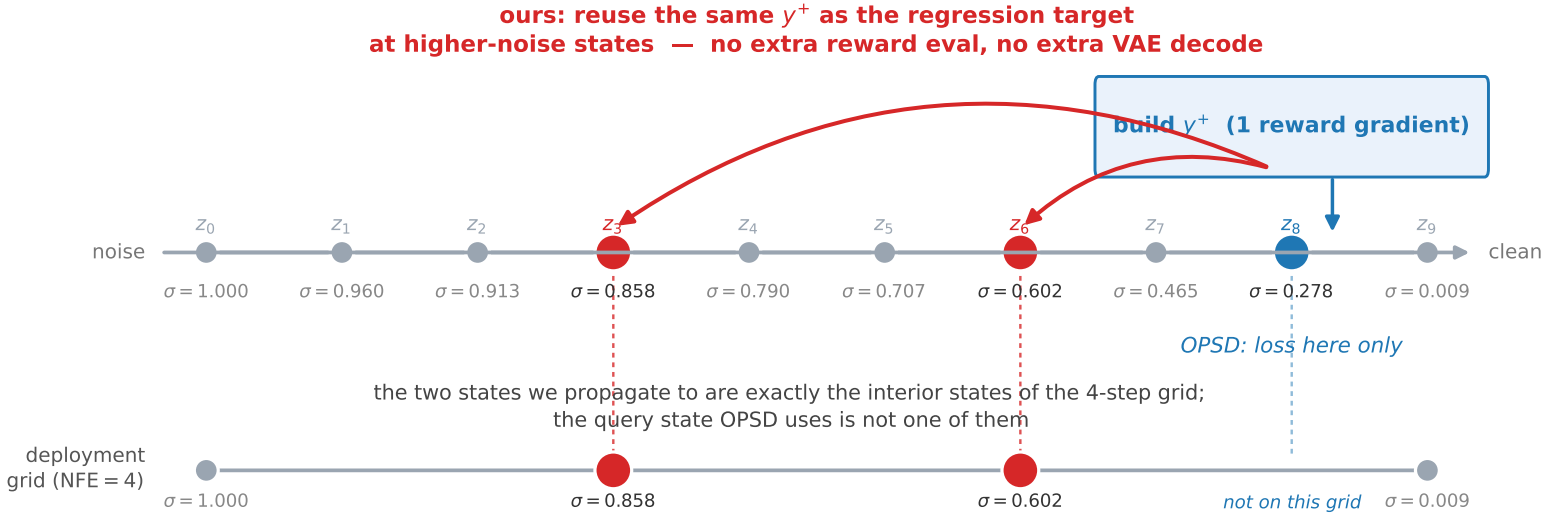
ours: reuse the same y^+ as the regression target
at higher-noise states — no extra reward eval, no extra VAE decode
build y^+ (1 reward gradient)
noise
clean
z_0
z_1
z_2
z_3
z_4
z_5
z_6
z_7
z_8
z_9
$\sigma = 1.000$
$\sigma = 0.960$
$\sigma = 0.913$
$\sigma = 0.858$
$\sigma = 0.790$
$\sigma = 0.707$
$\sigma = 0.602$
$\sigma = 0.465$
$\sigma = 0.278$
$\sigma = 0.009$
OPSD: loss here only
the two states we propagate to are exactly the interior states of the 4-step grid;
the query state OPSD uses is not one of them
deployment
grid (NFE = 4)
$\sigma = 1.000$
$\sigma = 0.858$
$\sigma = 0.602$
not on this grid
$\sigma = 0.009$