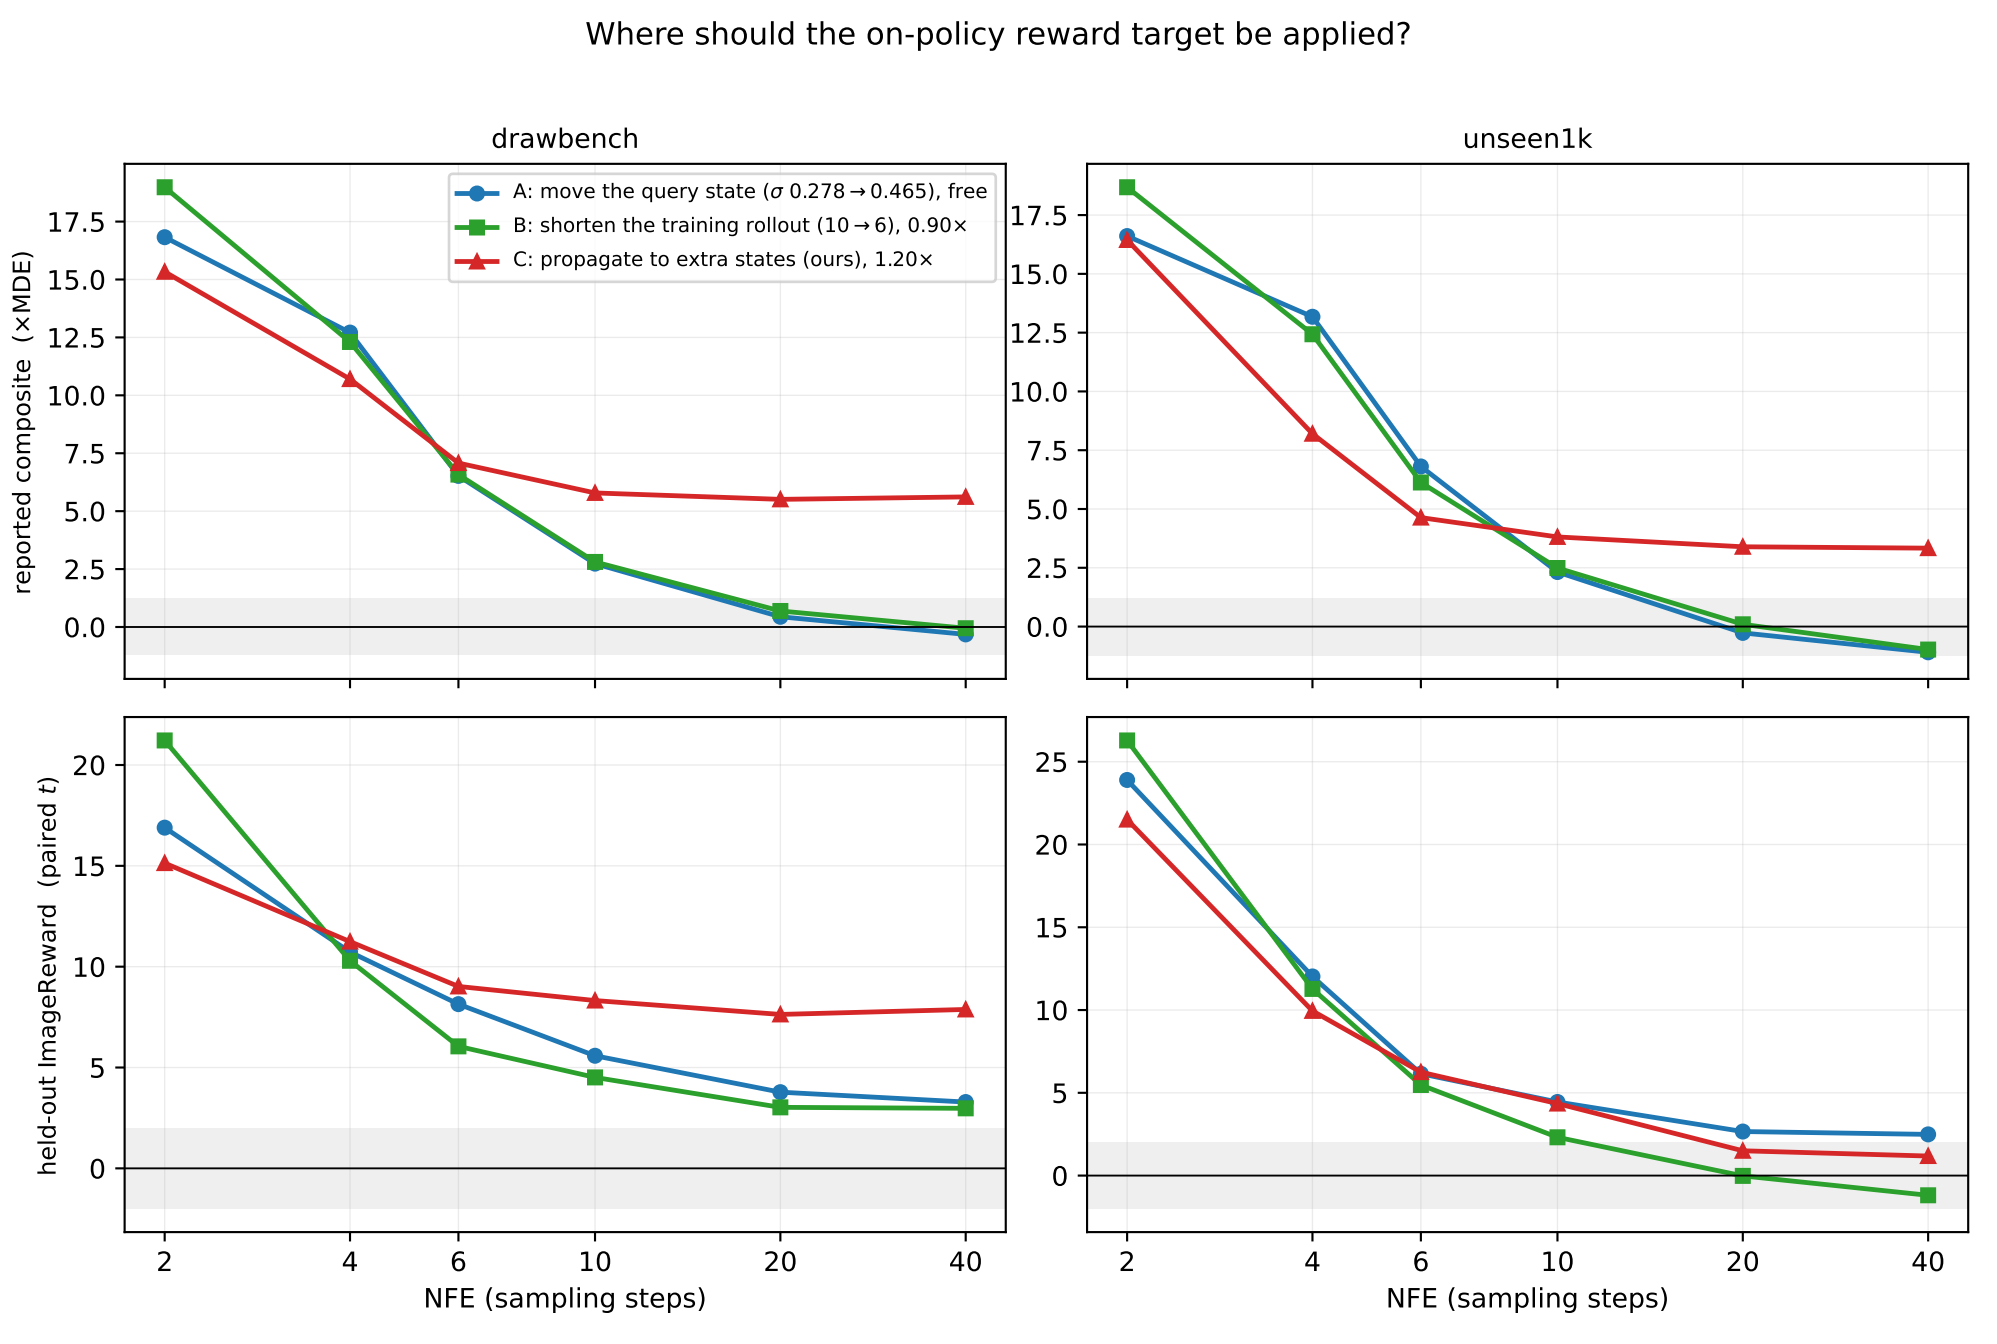
Where should the on-policy reward target be applied?
drawbench
unseen1k
A: move the query state (σ 0.278 → 0.465), free
B: shorten the training rollout (10 → 6), 0.90×
C: propagate to extra states (ours), 1.20×
reported composite (×MDE)
held-out ImageReward (paired t)
NFE (sampling steps)
NFE (sampling steps)