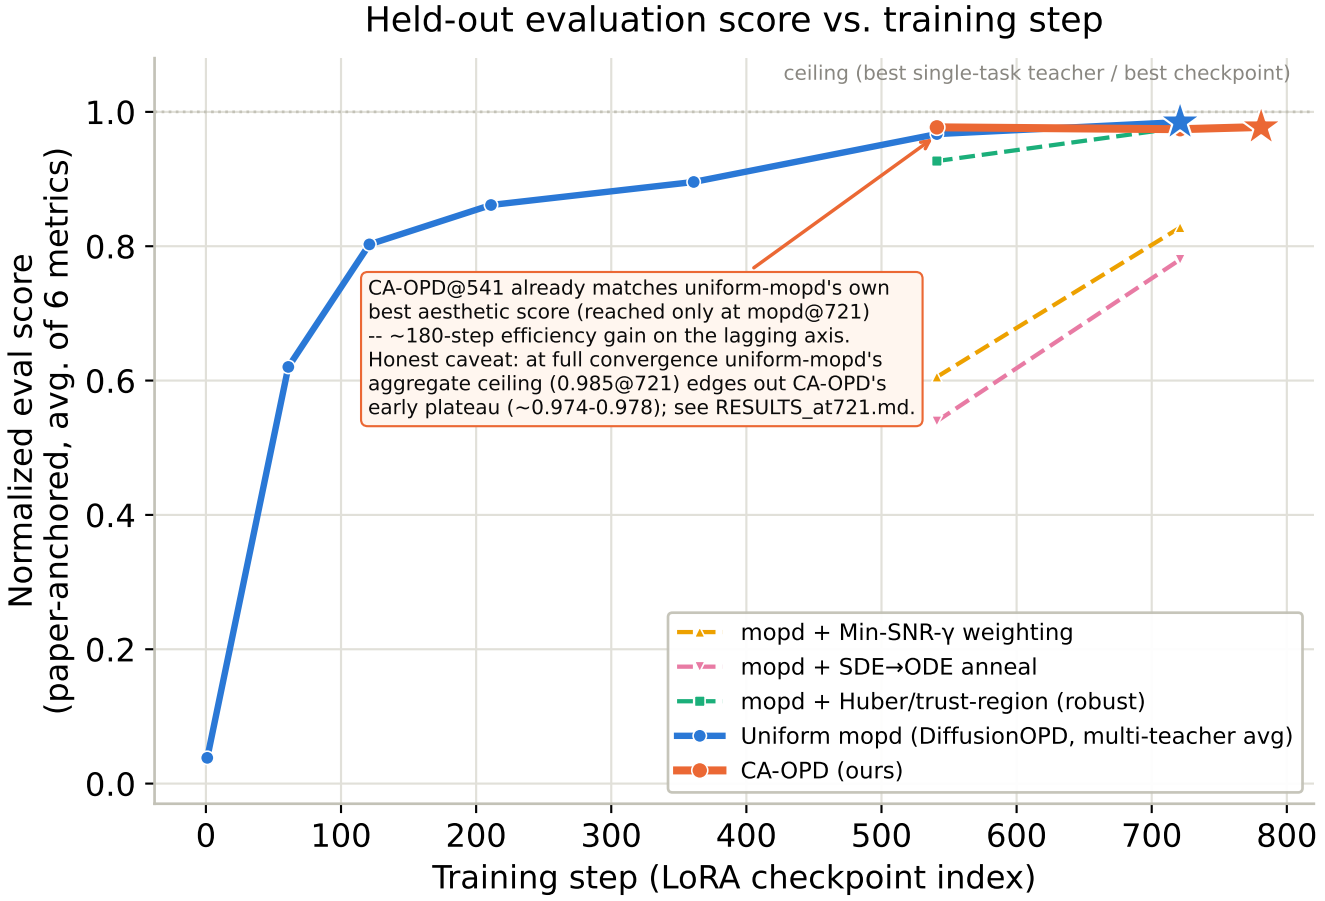
Held-out evaluation score vs. training step
ceiling (best single-task teacher / best checkpoint)
CA-OPD@541 already matches uniform-mopd's own best aesthetic score (reached only at mopd@721) -- ~180-step efficiency gain on the lagging axis. Honest caveat: at full convergence uniform-mopd's aggregate ceiling (0.985@721) edges out CA-OPD's early plateau (~0.974-0.978); see RESULTS_at721.md.
mopd + Min-SNR-γ weighting
mopd + SDE→ODE anneal
mopd + Huber/trust-region (robust)
Uniform mopd (DiffusionOPD, multi-teacher avg)
CA-OPD (ours)
Normalized eval score (paper-anchored, avg. of 6 metrics)
Training step (LoRA checkpoint index)
0.0
0.2
0.4
0.6
0.8
1.0
0
100
200
300
400
500
600
700
800